

MATH412/COMPSCI434/MATH713
Fall 2025

Topological Data Analysis

Topic 11: TDA + Machine Learning

Instructor: Ling Zhou

Machine Learning

Neural network

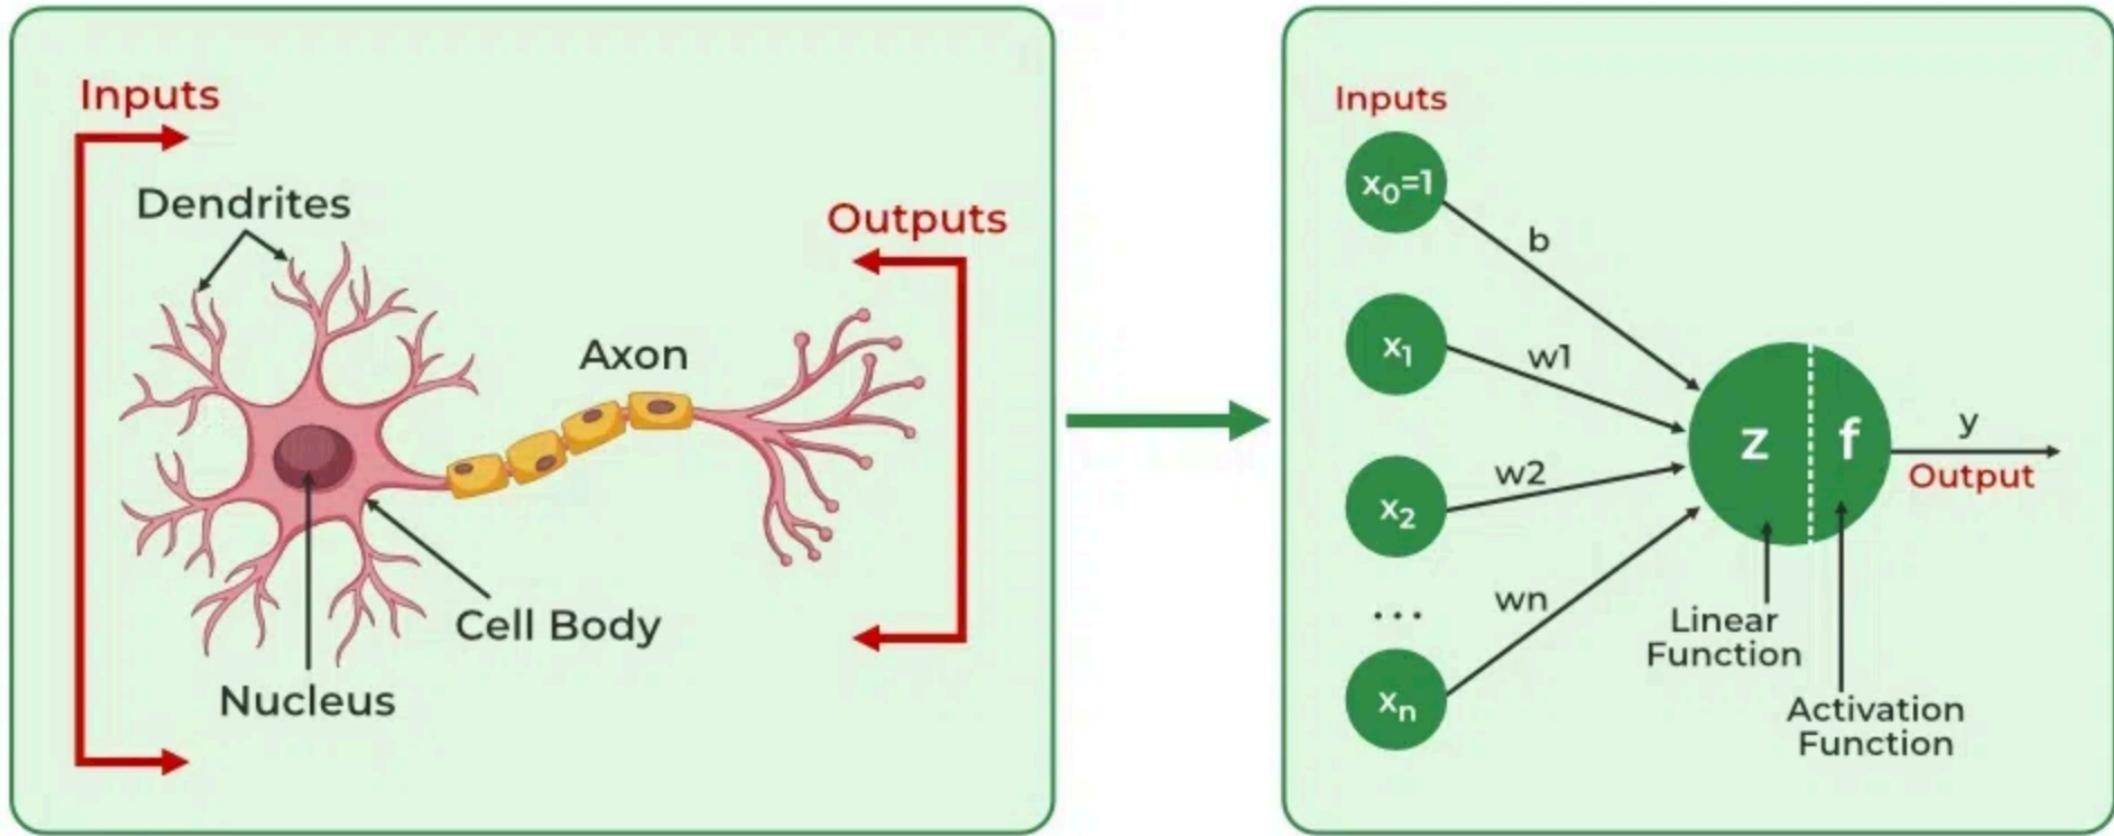
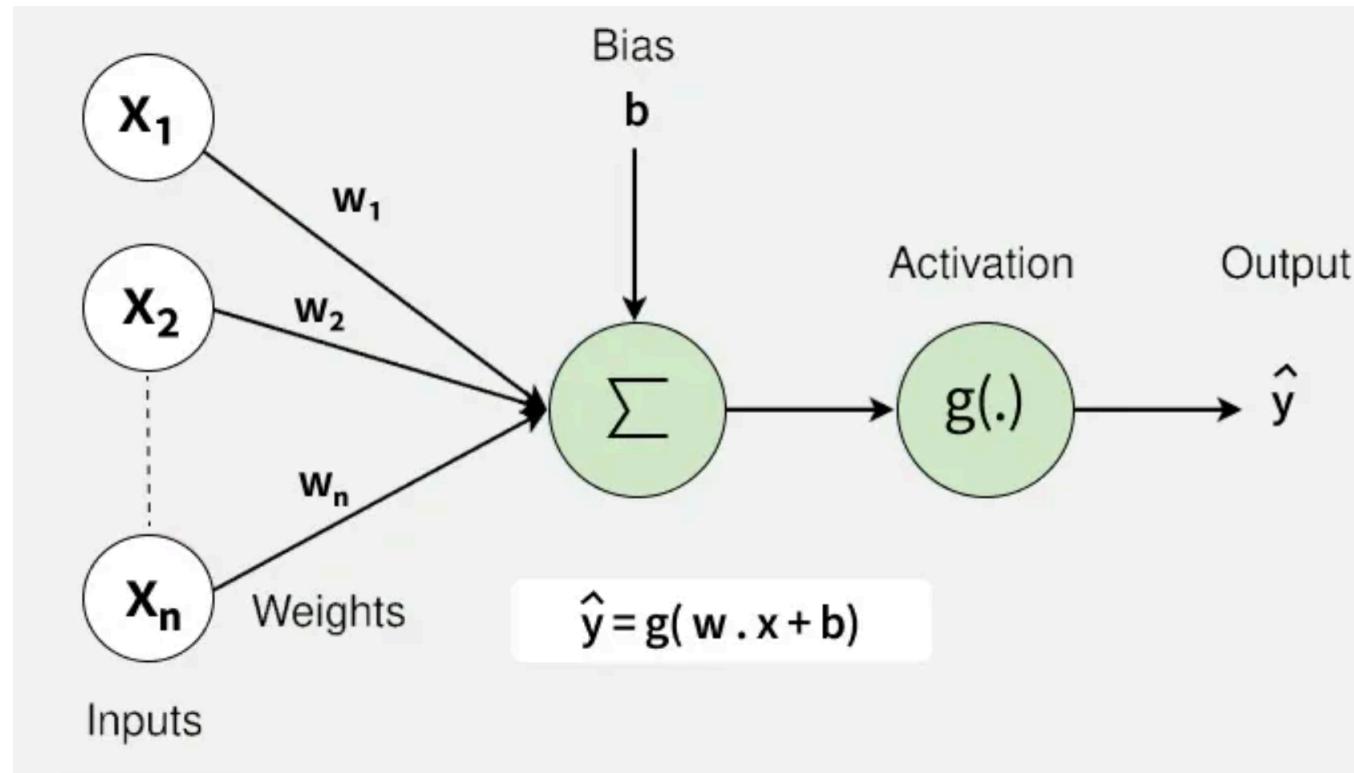


Figure taken from <https://www.geeksforgeeks.org/machine-learning/neural-networks-a-beginners-guide/>

Neural network: Architecture



$$\hat{y} = g(wx + b)$$

- x & b : vectors
- W : matrix
- g : activation function

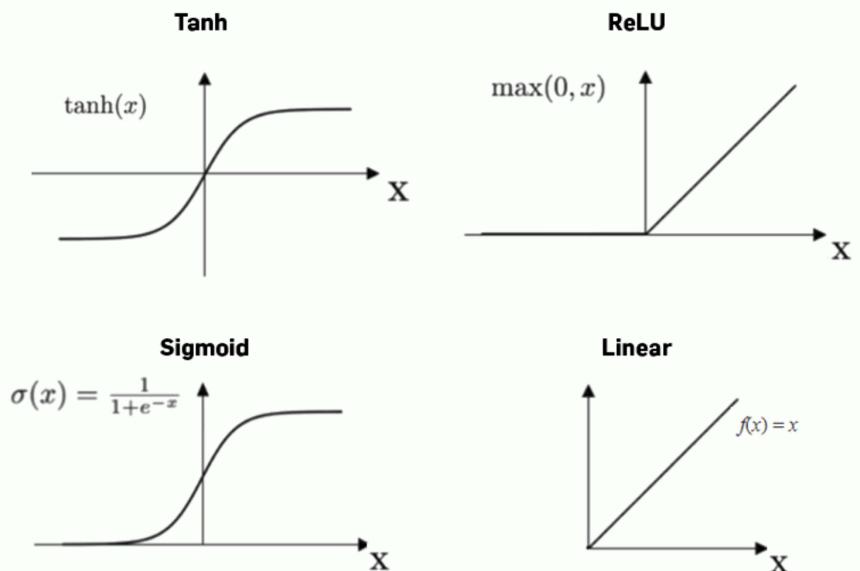


Figure taken from <https://www.geeksforgeeks.org/machine-learning/neural-networks-a-beginners-guide/>

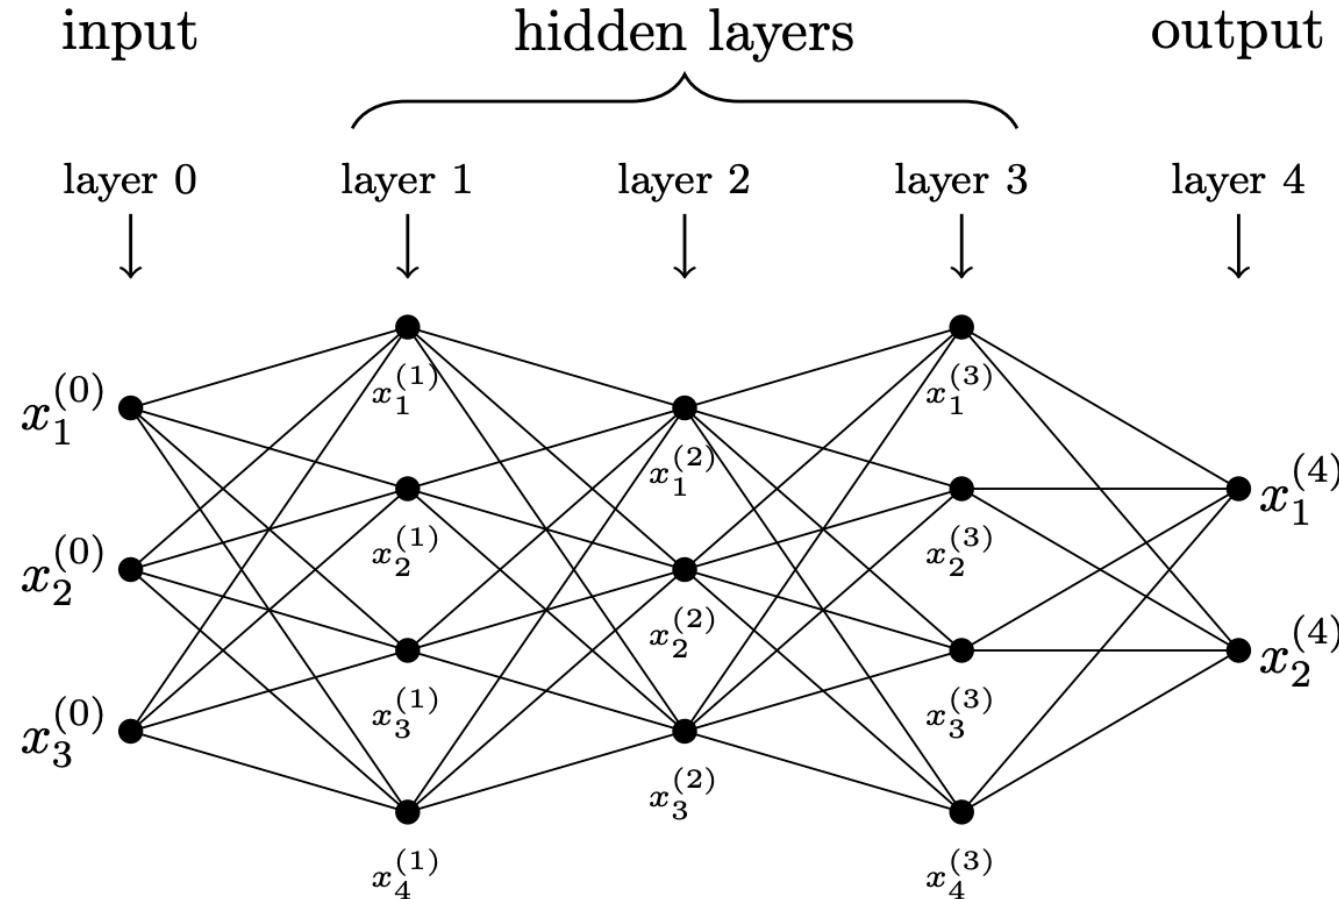
Neural network: Architecture

Philipp Petersen¹ and Jakob Zech²

¹Universität Wien, Fakultät für Mathematik, 1090 Wien, Austria,
philipp.petersen@univie.ac.at

²Universität Heidelberg, Interdisziplinäres Zentrum für Wissenschaftliches Rechnen, 69120
Heidelberg, Germany, jakob.zech@uni-heidelberg.de

April 8, 2025



Neural network: Architecture

Philipp Petersen¹ and Jakob Zech²

¹Universität Wien, Fakultät für Mathematik, 1090 Wien, Austria,
philipp.petersen@univie.ac.at

²Universität Heidelberg, Interdisziplinäres Zentrum für Wissenschaftliches Rechnen, 69120
Heidelberg, Germany, jakob.zech@uni-heidelberg.de

April 8, 2025

Definition 2.1. Let $L \in \mathbb{N}$, $d_0, \dots, d_{L+1} \in \mathbb{N}$, and let $\sigma: \mathbb{R} \rightarrow \mathbb{R}$. A function $\Phi: \mathbb{R}^{d_0} \rightarrow \mathbb{R}^{d_{L+1}}$ is called a **neural network** if there exist matrices $\mathbf{W}^{(\ell)} \in \mathbb{R}^{d_{\ell+1} \times d_\ell}$ and vectors $\mathbf{b}^{(\ell)} \in \mathbb{R}^{d_{\ell+1}}$, $\ell = 0, \dots, L$, such that with

$$\mathbf{x}^{(0)} := \mathbf{x} \tag{2.1.1a}$$

$$\mathbf{x}^{(\ell)} := \sigma(\mathbf{W}^{(\ell-1)} \mathbf{x}^{(\ell-1)} + \mathbf{b}^{(\ell-1)}) \quad \text{for } \ell \in \{1, \dots, L\} \tag{2.1.1b}$$

$$\mathbf{x}^{(L+1)} := \mathbf{W}^{(L)} \mathbf{x}^{(L)} + \mathbf{b}^{(L)} \tag{2.1.1c}$$

holds

$$\Phi(\mathbf{x}) = \mathbf{x}^{(L+1)} \quad \text{for all } \mathbf{x} \in \mathbb{R}^{d_0}.$$

We call L the **depth**, $d_{\max} = \max_{\ell=1, \dots, L} d_\ell$ the **width**, σ the **activation function**, and $(\sigma; d_0, \dots, d_{L+1})$ the **architecture** of the neural network Φ . Moreover, $\mathbf{W}^{(\ell)} \in \mathbb{R}^{d_{\ell+1} \times d_\ell}$ are the **weight matrices** and $\mathbf{b}^{(\ell)} \in \mathbb{R}^{d_{\ell+1}}$ the **bias vectors** of Φ for $\ell = 0, \dots, L$.

Neural network: Architecture

- ▶ This is the **Feedforward NN**. There are other types of NNs.

$$f^{(l)}(x) = \sigma(W^{(l)}x + b^{(l)}) \quad f^{(L)} \circ f^{(L-1)} \circ \dots \circ f^{(1)}(x)$$

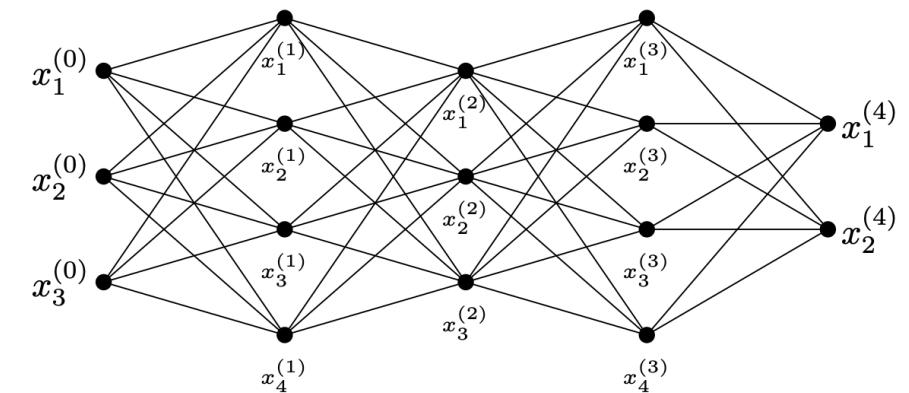
- ▶ **Recurrent NN**: information can flow backward.

$$f(x_t, h_{t-1}) = \sigma(W_{xh}x_t + W_{hh}h_{t-1} + b)$$

- ▶ **Convolutional NN**: replace matrix product with convolution

$$\text{ConvLayer}(x) = \sigma(w * x + b) \quad (x * w)(t) = \sum_{k=-\infty}^{\infty} x(k) w(t-k)$$

- ▶ **Residual NN**: one may skip connections.
- ▶ **Graph NN**: information flows following a graph
- ▶ **Transformer** (used in large language models):
 - ▶ A layer is self-attention + FNN (+ something else)
- ▶ etc



continuous functions on \mathbb{R}^d

Why neural networks work?

Definition 3.2. Let $d \in \mathbb{N}$. A set of functions \mathcal{H} from \mathbb{R}^d to \mathbb{R} is a **universal approximator** (of $C^0(\mathbb{R}^d)$), if for every $\varepsilon > 0$, every compact $K \subseteq \mathbb{R}^d$, and every $f \in C^0(\mathbb{R}^d)$, there exists $g \in \mathcal{H}$ such that $\sup_{x \in K} |f(x) - g(x)| < \varepsilon$. *Every cont. func. can be appr. by some func. in \mathcal{H}*

Definition 3.6. Let $d, m, L, n \in \mathbb{N}$ and $\sigma: \mathbb{R} \rightarrow \mathbb{R}$. The set of all functions realized by neural networks with d -dimensional input, m -dimensional output, depth at most L , width at most n , and activation function σ is denoted by

$$\mathcal{N}_d^m(\sigma; L, n) := \{\Phi: \mathbb{R}^d \rightarrow \mathbb{R}^m \mid \Phi \text{ as in Def. 2.1, } \text{depth}(\Phi) \leq L, \text{ width}(\Phi) \leq n\}.$$

Furthermore,

$$\mathcal{N}_d^m(\sigma; L) := \bigcup_{n \in \mathbb{N}} \mathcal{N}_d^m(\sigma; L, n).$$

\iff {FNN functions}

Theorem 3.8. Let $d \in \mathbb{N}$ and $\sigma \in \underline{\mathcal{M}}$. Then $\mathcal{N}_d^1(\sigma; 1)$ is a universal approximator of $C^0(\mathbb{R}^d)$ if and only if σ is not a polynomial. *\hookrightarrow a certain class of activation functions*

Neural network: Training

$$\theta = \{W^{(1)}, b^{(1)}, W^{(2)}, b^{(2)}, \dots, W^{(L)}, b^{(L)}\}$$

- ▶ Training a neural network means optimizing parameters to minimize loss:

$$\min_{\theta} \mathcal{L}(\theta)$$

- ▶ Supervised learning has labeled training sets (e.g. classification tasks):

$$\mathcal{L}_{\text{sup}}(\theta) = \frac{1}{N} \sum_i \ell(f_{\theta}(x_i), y_i)$$

- ▶ Unsupervised learning has just the data but no labels:

$$\mathcal{L}_{\text{unsup}}(\theta) = \text{some measure of how well the model represents } x_i$$

- ▶ Gradient descent: compute loss -> compute gradient -> update parameters

TDA in ML

- ▶ TDA can provide features, which can be used as inputs to NN
- ▶ TDA can be used to define topology-aware loss functions

Topological features as inputs to
NN

Deep learning with topological signatures

Deep Learning with Topological Signatures

Christoph Hofer
Department of Computer Science
University of Salzburg, Austria
chofer@cosy.sbg.ac.at

Roland Kwitt
Department of Computer Science
University of Salzburg, Austria
Roland.Kwitt@sbg.ac.at

Marc Niethammer
UNC Chapel Hill, NC, USA
mn@cs.unc.edu

Andreas Uhl
Department of Computer Science
University of Salzburg, Austria
uhl@cosy.sbg.ac.at

2017

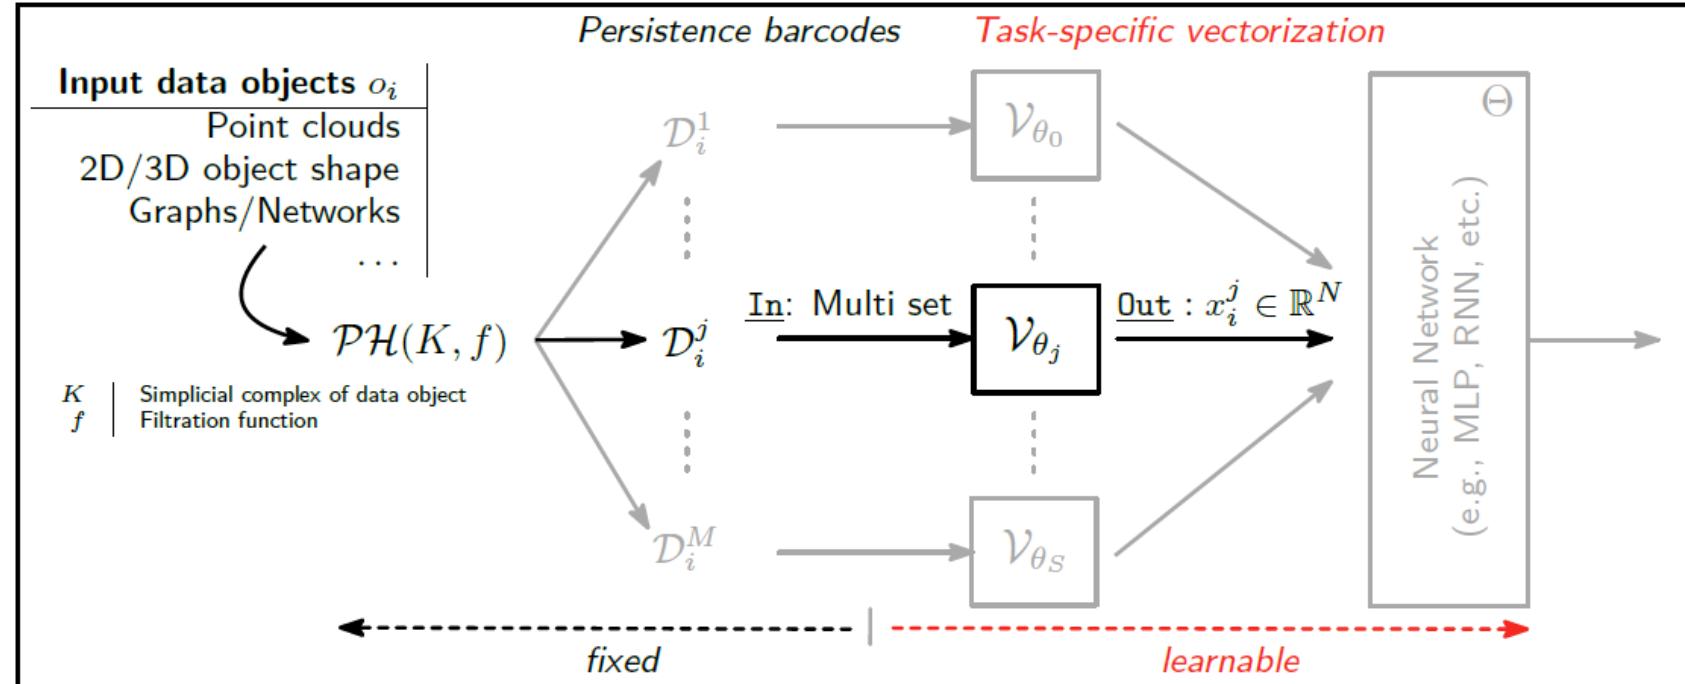
Abstract

Inferring topological and geometrical information from data can offer an alternative perspective on machine learning problems. Methods from topological data analysis, e.g., persistent homology, enable us to obtain such information, typically in the form of summary representations of topological features. However, such topological signatures often come with an unusual structure (e.g., multisets of intervals) that is highly impractical for most machine learning techniques. While many strategies have been proposed to map these topological signatures into machine learning compatible representations, they suffer from being agnostic to the target learning task. In contrast, we propose a technique that enables us to input topological signatures to deep neural networks and learn a task-optimal representation during training. Our approach is realized as a novel input layer with favorable theoretical properties. Classification experiments on 2D object shapes and social network graphs demonstrate the versatility of the approach and, in case of the latter, we even outperform the state-of-the-art by a large margin.

1 Introduction

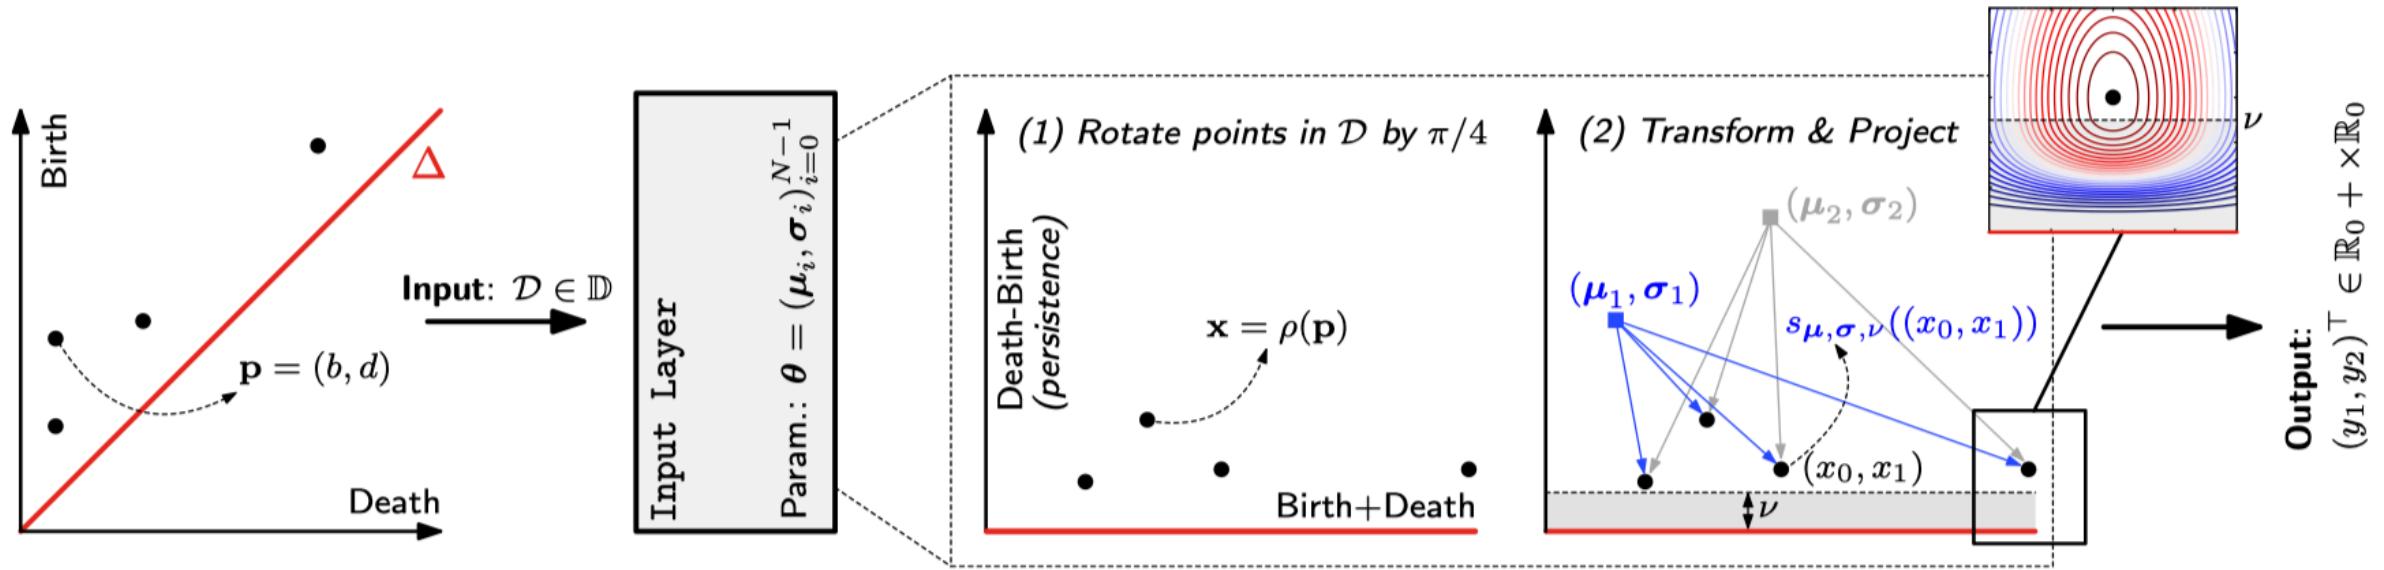
Methods from algebraic topology have only recently emerged in the machine learning community, most prominently under the term *topological data analysis (TDA)* [7]. Since TDA enables us to infer relevant topological and geometrical information from data, it can offer a novel and potentially beneficial perspective on various machine learning problems. Two compelling benefits of TDA are (1) its versatility, i.e., we are not restricted to any particular kind of data (such as images, sensor measurements, time-series, graphs, etc.) and (2) its robustness to noise. Several works have demonstrated that TDA can be beneficial in a diverse set of problems, such as studying the manifold of natural image patches [8], analyzing activity patterns of the visual cortex [28], classification of 3D surface meshes [27, 22], clustering [11], or recognition of 2D object shapes [29].

Currently, the most widely-used tool from TDA is *persistent homology* [15, 14]. Essentially¹, persistent homology allows us to track topological changes as we analyze data at multiple “scales”. As the scale changes, topological features (such as connected components, holes, etc.) appear and disappear. Persistent homology associates a *lifespan* to these features in the form of a *birth* and a *death* time. The collection of (birth, death) tuples forms a multiset that can be visualized as a persistence diagram or a barcode, also referred to as a *topological signature* of the data. However, leveraging these signatures for learning purposes poses considerable challenges, mostly due to their



¹We will make these concepts more concrete in Sec. 2.

Zoom in to the topological feature extraction



- ▶ The benefit of defining this new function s (instead of using other vectorization methods) is that s has *learnable* parameters

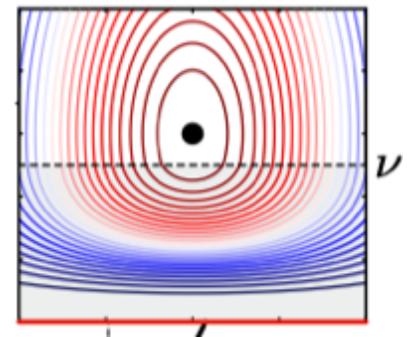
Zoom in to the topological feature extraction

Definition 3. Let $\mu = (\mu_0, \mu_1)^\top \in \mathbb{R} \times \mathbb{R}^+$, $\sigma = (\sigma_0, \sigma_1) \in \mathbb{R}^+ \times \mathbb{R}^+$ and $\nu \in \mathbb{R}^+$. We define

$$s_{\mu, \sigma, \nu} : \mathbb{R} \times \mathbb{R}_0^+ \rightarrow \mathbb{R}$$

as follows:

$$s_{\mu, \sigma, \nu}((x_0, x_1)) = \begin{cases} e^{-\sigma_0^2(x_0 - \mu_0)^2 - \sigma_1^2(x_1 - \mu_1)^2}, & x_1 \in [\nu, \infty) \\ e^{-\sigma_0^2(x_0 - \mu_0)^2 - \sigma_1^2(\ln(\frac{x_1}{\nu}) + \nu - \mu_1)^2}, & x_1 \in (0, \nu) \\ 0, & x_1 = 0 \end{cases}$$



A persistence diagram \mathcal{D} is then projected w.r.t. $s_{\mu, \sigma, \nu}$ via

$$S_{\mu, \sigma, \nu} : \mathbb{D} \rightarrow \mathbb{R}, \quad \mathcal{D} \mapsto \sum_{\mathbf{x} \in \mathcal{D}} s_{\mu, \sigma, \nu}(\rho(\mathbf{x})) . \quad (4)$$

Experimental results

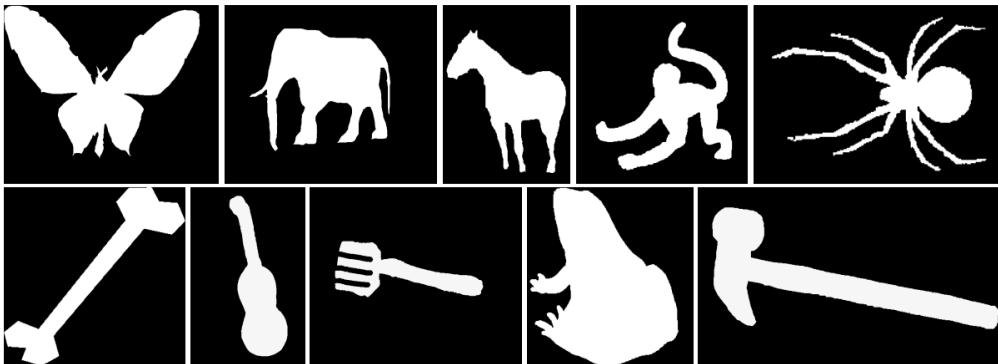


Figure 3: *Left:* some examples from the MPEG-7 (*bottom*) and Animal (*top*) datasets. *Right:* Classification results, compared to the two best (\dagger) and two worst (\ddagger) results reported in [30].

- Advantages: no preprocessing required; easily generalizable to 3D images

	5	10	20	40	80	160	Ours
MPEG-7	81.8	82.3	79.7	74.5	68.2	64.4	91.8
Animal	48.8	50.0	46.2	42.4	39.3	36.0	69.5
reddit-5k	37.1	38.2	39.7	42.1	43.8	45.2	54.5
reddit-12k	24.2	24.6	27.9	29.8	31.5	31.6	44.5

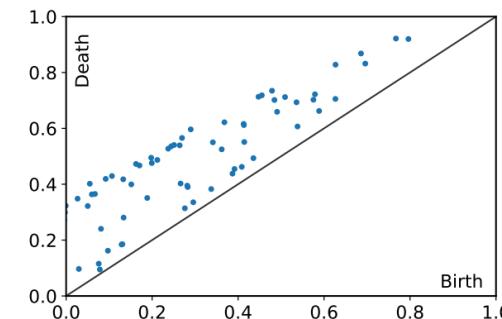
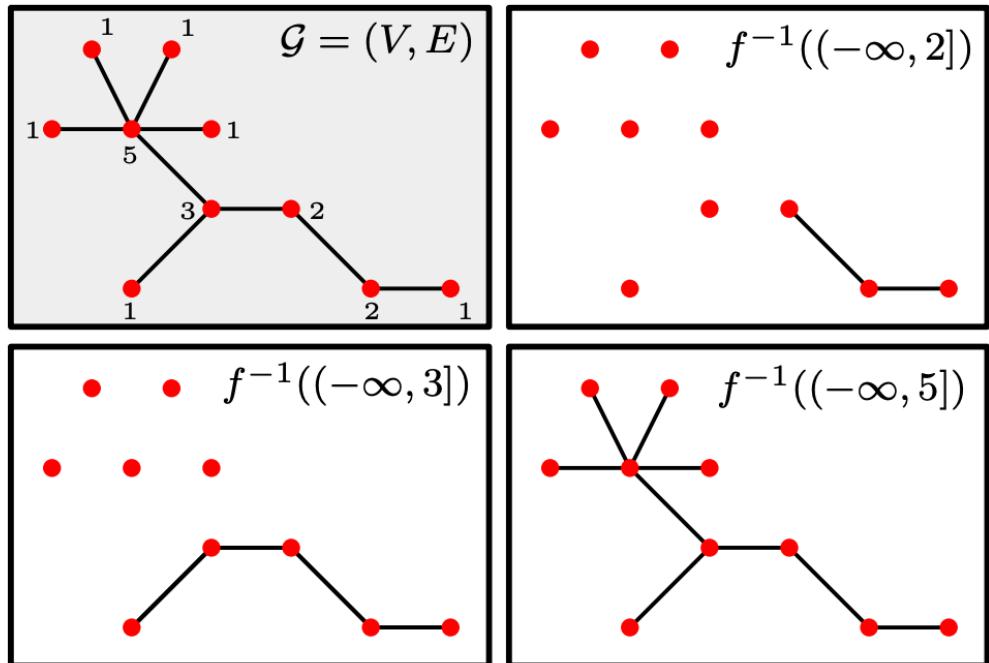


Figure 4: *Left:* Classification accuracies for a linear SVM trained on vectorized (in \mathbb{R}^N) persistence diagrams (see Sec. 5.3). *Right:* Exemplary visualization of the learned structure elements (in 0-th dimension) for the Animal dataset and filtration direction $\mathbf{d} = (-1, 0)^\top$. Centers of the learned elements are marked in blue.

Experimental results

► Graph classification on social network data



	reddit-5k	reddit-12k
GK [31]	41.0	31.8
DGK [31]	41.3	32.2
PSCN [24]	49.1	41.3
RF [4]	50.9	42.7
Ours (w/o essential)	49.1	38.5
Ours (w/ essential)	54.5	44.5

Figure 5: *Left:* Illustration of graph filtration by vertex degree, i.e., $f \equiv \deg$ (for different choices of a_i , see Sec. 2). *Right:* Classification results as reported in [31] for GK and DGK, Patchy-SAN (PSCN) as reported in [24] and feature-based random-forest (RF) classification from [4].

PersLay

- ▶ [Carriere et al., 2020]

- ▶ Given a persistence diagram $D \in \mathfrak{D}$,

$$\text{Perslay}(D) := \text{op}(\left\{ \omega(p) \cdot \phi(p) \right\}_{p \in D})$$

- ▶ where $\omega: R^2 \rightarrow R$ is a weight function for points in persistence diagram, and
- ▶ $\phi: R^2 \rightarrow R^k$ is a point transformation function to map each persistent point to a k -vector,
- ▶ and **op** is a permutation invariant set function
 - e.g, max, sum, k-th largest value, etc

PersLay: A Neural Network Layer for Persistence Diagrams and New Graph Topological Signatures

Mathieu Carrière
Columbia University

Frédéric Chazal
Inria Saclay

Yuichi Ike
Fujitsu Laboratories

Théo Lacombe
Inria Saclay

Martin Royer
Inria Saclay

Yuhei Umeda
Fujitsu Laboratories

Abstract

Persistence diagrams, the most common descriptors of Topological Data Analysis, encode topological properties of data and have already proved pivotal in many different applications of data science. However, since the metric space of persistence diagrams is not Hilbert, they end up being difficult inputs for most Machine Learning techniques. To address this concern, several vectorization methods have been put forward that embed persistence diagrams into either finite-dimensional Euclidean space or implicit infinite dimensional Hilbert space with kernels.

In this work, we focus on persistence diagrams built on top of graphs. Relying on extended persistence theory and the so-called heat kernel signature, we show how graphs can be encoded by (extended) persistence diagrams in a provably stable way. We then propose a general and versatile framework for learning vectorizations of persistence diagrams, which encompasses most of the vectorization techniques used in the literature. We finally showcase the experimental strength of our setup by achieving competitive scores on classification tasks on real-life graph datasets.

1 INTRODUCTION

Topological Data Analysis (TDA) is a field of data science whose goal is to detect and encode topological

Proceedings of the 23rd International Conference on Artificial Intelligence and Statistics (AISTATS) 2020, Palermo, Italy. PMLR: Volume 108. Copyright 2020 by the author(s).

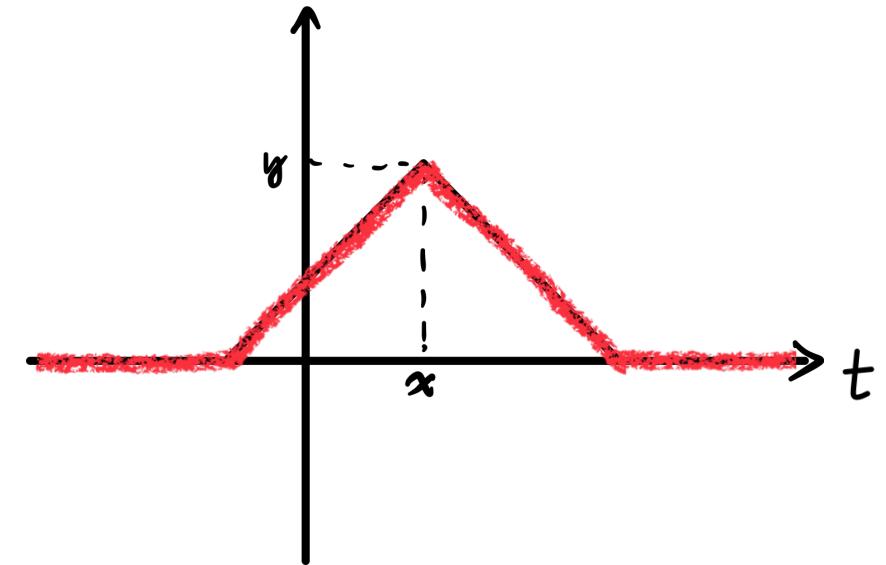
features (such as connected components, loops, cavities...) that are present in datasets in order to improve inference and prediction. Its main descriptor is the so-called *persistence diagram*, which takes the form of a set of points in the Euclidean plane R^2 , each point corresponding to a topological feature of the data, with its coordinates encoding the feature size. This descriptor has been successfully used in many different applications of data science, such as signal analysis (Perea & Harer, 2015), material science (Buchet et al., 2018), cellular data (Cámarra, 2017), or shape recognition (Li et al., 2014) to name a few. This wide range of applications is mainly due to the fact that persistence diagrams encode information based on topology, and as such this information is very often complementary to the one retrieved by more classical descriptors.

However, the space of persistence diagrams heavily lacks structure: different persistence diagrams may have different number of points, and several basic operations are not well-defined, such as addition and scalar multiplication, which unfortunately dramatically impedes their use in machine learning applications. To handle this issue, a lot of attention has been devoted to *vectorizations* of persistence diagrams through the construction of either *finite-dimensional embeddings* (Adams et al., 2017; Carrière et al., 2015; Chazal et al., 2015; Kalisnik, 2018), i.e., embeddings turning persistence diagrams into vectors in Euclidean space R^d , or *kernels* (Bubenik, 2015; Carrière et al., 2017; Kusano et al., 2016; Le & Yamada, 2018; Reininghaus et al., 2015), i.e., generalized scalar products that implicitly turn persistence diagrams into elements of infinite-dimensional Hilbert spaces.

Even though these methods improved the use of persistence diagrams in machine learning tremendously, several issues remain. For instance, most of these vectorizations only have a few trainable parameters, which may prevent them from fitting well to specific appli-

Examples of ϕ

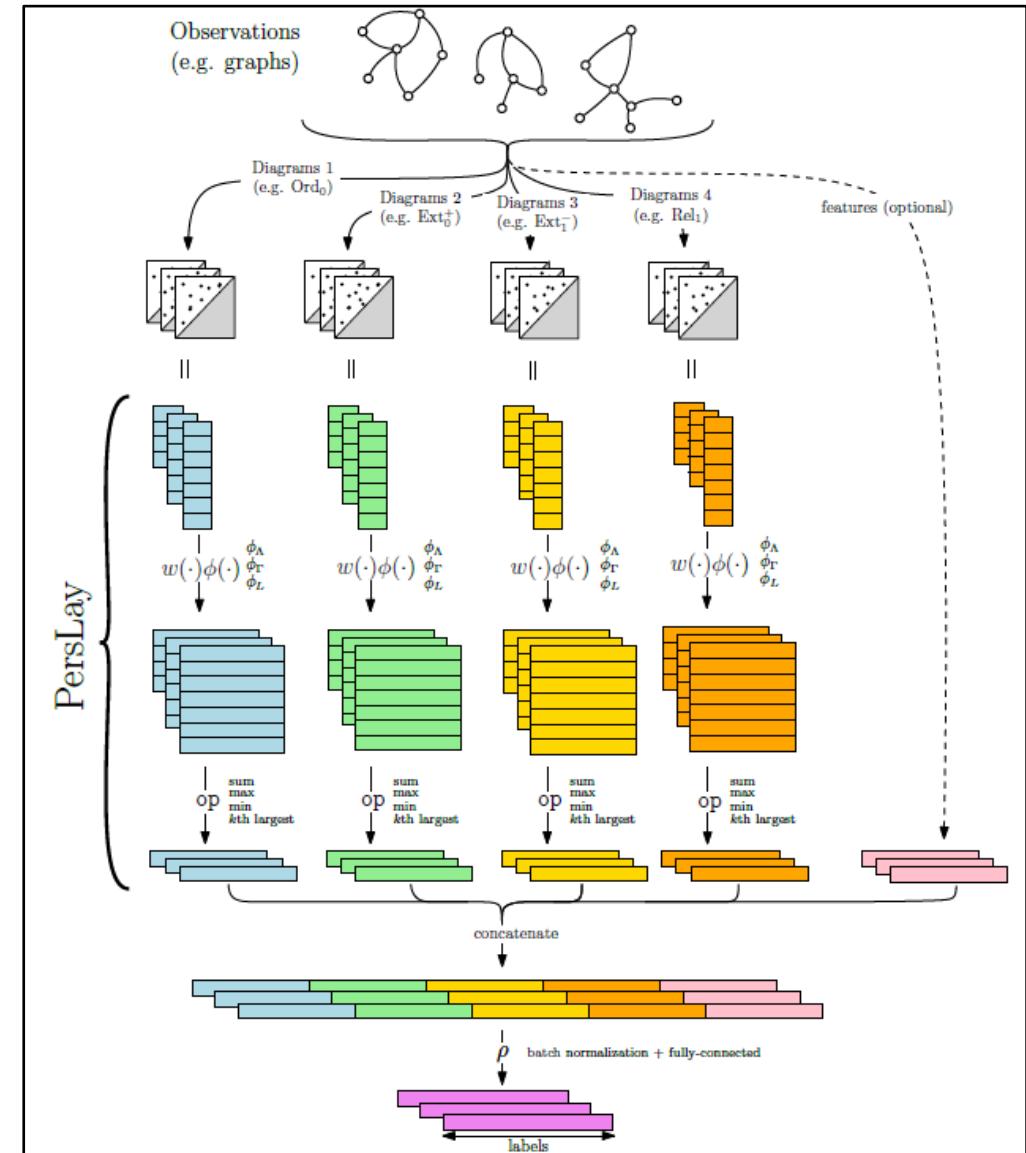
- The *triangle point transformation* $\phi_\Lambda : \mathbb{R}^2 \rightarrow \mathbb{R}^q, p \mapsto [\Lambda_p(t_1), \Lambda_p(t_2), \dots, \Lambda_p(t_q)]^T$ where the triangle function Λ_p associated to a point $p = (x, y) \in \mathbb{R}^2$ is $\Lambda_p : t \mapsto \max\{0, y - |t - x|\}$, with $q \in \mathbb{N}$ and $t_1, \dots, t_q \in \mathbb{R}$.
- The *Gaussian point transformation* $\phi_\Gamma : \mathbb{R}^2 \rightarrow \mathbb{R}^q, p \mapsto [\Gamma_p(t_1), \Gamma_p(t_2), \dots, \Gamma_p(t_q)]^T$, where the Gaussian function Γ_p associated to a point $p = (x, y) \in \mathbb{R}^2$ is $\Gamma_p : t \mapsto \exp(-\|p - t\|_2^2/(2\sigma^2))$ for a given $\sigma > 0$, $q \in \mathbb{N}$ and $t_1, \dots, t_q \in \mathbb{R}^2$.
- The *line point transformation* $\phi_L : \mathbb{R}^2 \rightarrow \mathbb{R}^q, p \mapsto [L_{\Delta_1}(p), L_{\Delta_2}(p), \dots, L_{\Delta_q}(p)]^T$, where the line function L_Δ associated to a line Δ with direction vector $e_\Delta \in \mathbb{R}^2$ and bias $b_\Delta \in \mathbb{R}$ is $L_\Delta : p \mapsto \langle p, e_\Delta \rangle + b_\Delta$, with $q \in \mathbb{N}$ and $\Delta_1, \dots, \Delta_q$ are q lines in the plane.



- Using: $\phi = \phi_\Lambda$ with samples $t_1, \dots, t_q \in \mathbb{R}$, $\text{op} = k$ th largest value, $w = 1$ (a constant weight function), amounts to evaluating the *kth persistence landscape* (Bubenik, 2015) on $t_1, \dots, t_q \in \mathbb{R}$.

PersLay recovers other vectorizations

- ▶ PersLay is a general framework
- ▶ By choosing specific ω and ϕ , can recover most other feature vectorization,
 - ▶ persistence landscapes
 - ▶ persistence images
 - ▶ Hofer's network structure
- ▶ Have better performance than previous kernel methods; comparable performance against GNNs



Learning metrics for persistence-based summaries and applications for graph classification

Qi Zhao
zhao.2017@osu.edu

Yusu Wang
yusu@cse.ohio-state.edu

Computer Science and Engineering Department
The Ohio State University
Columbus, OH 43221

2019

Abstract

Recently a new feature representation framework based on a topological tool called persistent homology (and its persistence diagram summary) has gained much momentum. A series of methods have been developed to map a persistence diagram to a vector representation so as to facilitate the downstream use of machine learning tools. In these approaches, the importance (weight) of different persistence features are usually *pre-set*. However often in practice, the choice of the weight-function should depend on the nature of the specific data at hand. It is thus highly desirable to *learn* a best weight-function (and thus metric for persistence diagrams) from labelled data. We study this problem and develop a new weighted kernel, called *WKPI*, for persistence summaries, as well as an optimization framework to learn the weight (and thus kernel). We apply the learned kernel to the challenging task of graph classification, and show that our WKPI-based classification framework obtains similar or (sometimes significantly) better results than **the best results** from a range of previous graph classification frameworks on benchmark datasets.

1 Introduction

In recent years a new data analysis methodology based on a topological tool called persistent homology has started to attract momentum. The persistent homology is one of the most important developments in the field of topological data analysis, and there have been fundamental developments both on the theoretical front (e.g., [23, 10, 13, 8, 14, 5]), and on algorithms / implementations (e.g., [43, 4, 15, 20, 29, 3]). On the high level, given a domain X with a function $f : X \rightarrow \mathbb{R}$ on it, the persistent homology summarizes “features” of X across multiple scales simultaneously in a single summary called the *persistence diagram* (see the second picture in Figure 1). A persistence diagram consists of a multiset of points in the plane, where each point $p = (b, d)$ intuitively corresponds to the birth-time (b) and death-time (d) of some (topological) features of X w.r.t. f . Hence it provides a concise representation of X , capturing *multi-scale features* of it simultaneously. Furthermore, the persistent homology framework can be applied to complex data (e.g. 3D shapes, or graphs), and different summaries could be constructed by putting different descriptor functions on input data.

Due to these reasons, a new persistence-based feature vectorization and data analysis framework (Figure 1) has become popular. Specifically, given a collection of objects, say a set of graphs modeling chemical compounds, one can first convert each shape to a persistence-based representation. The input data can now be viewed as a set of points in a persistence-based feature space. Equipping this space with appropriate distance or kernel, one can then perform downstream data analysis tasks (e.g., clustering).

- ▶ An extension of persistence image
- ▶ A learnable weight function
- ▶ Suitable for graph classification

Graph classification

▶ Benchmark datasets:

Datasets	graph #	class #	average_nodes #	average edges #	label #
MUTAG	188	2	17.93	19.79	7
PTC	344	2	14.29	14.69	19
ENZYME	600	6	32.63	64.14	3
PROTEIN	1113	2	39.06	72.82	3
DD	1178	2	284.32	715.66	81
NCI1	4110	2	29.87	32.30	37
IMDB BINARY	1000	2	19.77	96.53	-
IMDB MULTI	1500	3	13.00	65.94	-
REDDIT BINARY	2000	2	429.63	497.75	-
REDDIT 5K	4999	5	508.82	594.87	-
REDDIT 12K	12929	11	391.41	456.89	-

Statistics of benchmark graph data sets

Notes, some are attributed graphs.

Graph classification, cont.

- ▶ Comparison methods:
 - ▶ Weisfeiler-Lehman Kernel (WL) [Shervashidze et al., 2011]
 - ▶ Graphlet kernel (GK) [Shervashidze et al, 2009]
 - ▶ Deep Graph Kernel (DGK) [Yanardag and Vishwanathan, 2015]
 - ▶ FGSD [Verma and Zhang, 2017]
 - ▶ RetGK [Zhang et al., 2018]
- ▶ Graph isomorphism network (GIN) [Xu et al., 2018]
- ▶ PATCHYSAN (PSCN) [Niepert et al., 2016]

Graph classification, cont.

- ▶ Classification accuracy + variance
 - ▶ by SVM + a metric-learning approach by [Zhao, W., 2019]
 - ▶ 10 times 10-fold cross validation

Dataset	RetGK	WL	GK	DGK	PSCN	GIN	WKPI-kM	WKPI-kC
NCI1	84.5±0.2	85.4±0.3	62.3±0.3	80.3±0.5	76.3±1.7	82.7±1.6	87.2±0.4	84.7±0.4
NCI109	-	84.5±0.2	66.6±0.2	80.3±0.3	-	-	85.6±0.3	87.3±0.3
PTC	62.5±1.6	55.4±1.5	57.3±1.1	60.1±2.5	62.3±5.7	66.6±6.9	63.1±2.4	67.1±2.2
PROTEIN	75.8±0.6	71.2±0.8	71.7±0.6	75.7±0.5	75.0±2.5	76.2±2.6	78.8±0.4	74.9±0.3
DD	81.6±0.3	78.6±0.4	78.5±0.3	-	76.2±2.6	-	82.0±0.5	80.3±0.4
MUTAG	90.3±1.1	84.4±1.5	81.6±2.1	87.4±2.7	89.0±4.4	90.0±8.8	86.9±2.5	87.5±2.6
IMDB-BINARY	71.9±1.0	70.8±0.5	65.9±1.0	67.0±0.6	71.0±2.3	75.1±5.1	70.7±1.1	75.4±1.1
IMDB-MULTI	47.7±0.3	49.8±0.5	43.9±0.4	44.6±0.4	45.2±2.8	52.3 ±2.8	46.4±0.5	49.5±0.4
REDDIT-5K	56.1±0.5	51.2±0.3	41.0±0.2	41.3±0.2	49.1±0.7	57.5±1.5	58.5±0.4	60.2±0.6
REDDIT-12K	48.7±0.2	32.6±0.3	31.8±0.1	32.2±0.1	41.3±0.4	-	47.7±0.5	48.6±0.5

2019

Bastian Rieck ^{*} ¹ Christian Bock ^{*} ¹ Karsten Borgwardt ¹

Abstract

The Weisfeiler–Lehman graph kernel exhibits competitive performance in many graph classification tasks. However, its subtree features are not able to capture connected components and cycles, topological features known for characterising graphs. To extract such features, we leverage propagated node label information and transform unweighted graphs into metric ones. This permits us to augment the subtree features with topological information obtained using *persistent homology*, a concept from topological data analysis. Our method, which we formalise as a generalisation of Weisfeiler–Lehman subtree features, exhibits favourable classification accuracy and its improvements in predictive performance are mainly driven by including cycle information.

1. Introduction

Graph-structured data sets are ubiquitous in a variety of different application domains, each of them posing a separate challenge while also requiring different tasks to be solved. A common task involves *graph classification*, for which a variety of methods exists. These methods comprise convolutional neural networks (Duvenaud et al. 2015), recurrent neural networks (Lei et al. 2017), or Hilbert space methods (Vishwanathan et al. 2010), the latter also being referred to as *graph kernels*. While several approaches for defining graph kernels exist, the most common one uses the \mathcal{R} -convolution framework (Haussler 1999), which makes it possible to define the similarity between two graphs as a function of the similarity of their substructures.

Substructures that have been used for graph classification range from graphlets (Shervashidze et al. 2009), i.e. small non-isomorphic graphs of fixed size, over shortest

paths (Borgwardt & Kriegel 2005), to random walks (Gärtner et al. 2003, Kashima et al. 2003, Sugiyama & Borgwardt 2015). One of the most powerful substructures is the set of *subtree patterns* (Ramon & Gärtner 2003), i.e. patterns based on rooted subgraphs of a graph. Their computational complexity made their applicability somewhat limited, until the Weisfeiler–Lehman (WL) graph kernel framework (Shervashidze & Borgwardt 2009, Shervashidze et al. 2011) was developed. Properly trained, it still constitutes the state-of-the-art method for many graph classification tasks. The framework is based on the idea of iteratively propagating (node) label information through a graph, leading to a feature vector representation that can be used to assess the dissimilarity of two graphs.

One of the disadvantages of this framework is that its relabelling step, i.e. the step in which subtree patterns are being *compressed*, is somewhat “brittle”: labels are only compared with a Dirac kernel, making their dissimilarity a coarse function. Moreover, the subtree feature vector only contains counts of compressed labels and can neither account for their relevance with respect to the topology of the graph nor capture connected components and cycles, both of which are important and interpretable features for characterising graphs (Rieck et al. 2018, Sizemore et al. 2017). We thus propose an enhancement of the original WL stabilisation procedure that uses recent advances in topological data analysis (Munch 2017) to alleviate these issues. Our contributions are as follows:

- We measure the relevance of topological features (connected components and cycles) in graphs and use them to define a novel set of WL subtree features, which we show to be a generalised version of the original ones.
- We develop a topology-based kernel that uses an iterative variant of the WL stabilisation procedure to classify non-attributed graphs.
- We demonstrate that our proposed features perform favourably on a range of graph classification benchmark data sets. In particular, we empirically show that the inclusion of *cycle* information yields classification accuracy improvements over state-of-the-art methods.

^{*}Equal contribution ¹Department of Biosystems Science and Engineering, ETH Zurich, 4058 Basel, Switzerland. Correspondence to: Bastian Rieck <bastian.rieck@bsse.ethz.ch>, Karsten Borgwardt <karsten.borgwardt@bsse.ethz.ch>.

- ▶ Classic WL test provides a procedure to vectorize (labeled) graphs
- ▶ Persistent homology encodes topological info into the above procedure

Classification results

	D & D	MUTAG	NCI1	NCI109	PROTEINS	PTC-MR	PTC-FR	PTC-MM	PTC-FM
V-Hist	78.32 ± 0.35	85.96 ± 0.27	64.40 ± 0.07	63.25 ± 0.12	72.33 ± 0.32	58.31 ± 0.27	68.13 ± 0.23	66.96 ± 0.51	57.91 ± 0.83
E-Hist	72.90 ± 0.48	85.69 ± 0.46	63.66 ± 0.11	63.27 ± 0.07	72.14 ± 0.39	55.82 ± 0.00	65.53 ± 0.00	61.61 ± 0.00	59.03 ± 0.00
RetGK*	81.60 ± 0.30	90.30 ± 1.10	84.50 ± 0.20		75.80 ± 0.60	62.15 ± 1.60	67.80 ± 1.10	67.90 ± 1.40	63.90 ± 1.30
WL	79.45 ± 0.38	87.26 ± 1.42	85.58 ± 0.15	84.85 ± 0.19	76.11 ± 0.64	63.12 ± 1.44	67.64 ± 0.74	67.28 ± 0.97	64.80 ± 0.85
Deep-WL*		82.94 ± 2.68	80.31 ± 0.46	80.32 ± 0.33	75.68 ± 0.54	60.08 ± 2.55			
P-WL	79.34 ± 0.46	86.10 ± 1.37	85.34 ± 0.14	84.78 ± 0.15	75.31 ± 0.73	63.07 ± 1.68	67.30 ± 1.50	68.40 ± 1.17	64.47 ± 1.84
P-WL-C	78.66 ± 0.32	90.51 ± 1.34	85.46 ± 0.16	84.96 ± 0.34	75.27 ± 0.38	64.02 ± 0.82	67.15 ± 1.09	68.57 ± 1.76	65.78 ± 1.22
P-WL-UC	78.50 ± 0.41	85.17 ± 0.29	85.62 ± 0.27	85.11 ± 0.30	75.86 ± 0.78	63.46 ± 1.58	67.02 ± 1.29	68.01 ± 1.04	65.44 ± 1.18

Qi Zhao*

Ze Ye[†]

Chao Chen[†]

Yusu Wang*

Abstract

Local structural information can increase the adaptability of graph convolutional networks to large graphs with heterogeneous topology. Existing methods only use relatively simple topological information, such as node degrees. We present a novel approach leveraging advanced topological information, i.e., persistent homology, which measures the information flow efficiency at different parts of the graph. To fully exploit such structural information in real world graphs, we propose a new network architecture which learns to use persistent homology information to reweight messages passed between graph nodes during convolution. For node classification tasks, our network outperforms existing ones on a broad spectrum of graph benchmarks.

1 INTRODUCTION

Deep learning methods have achieved immense success in different domains such as computer vision and natural language processing (Goodfellow et al., 2016). While deep neural networks have shown strong performance on image or text data, their learning power is yet to be fully exploited on graph-structured data. At the same time, data with a latent graph structure is ubiquitous in modern data science. It is highly desirable to develop deep learning techniques that best suite graph structures, such as social network, knowledge network, brain connectivity network, etc.

Earlier works on Graph Neural Networks (GNNs) (Gori et al., 2005; Scarselli et al., 2009) use recursive networks. Those GNNs process the graph using a set of neurons, each corresponding to a node in the

graph. The neurons update nodes representation and exchange information from linked neighbor nodes iteratively until reaching equilibrium.

Inspired by the power of convolutional networks on image and text data, different ideas have been proposed to implement the “convolution” on graph structures. There are two main directions, spectral convolutions and spatial convolutions. Spectral convolutional networks (Bruna et al., 2014; Defferrard et al., 2016) apply convolutions to the spectral domain or the frequency domain of the input graph. These methods tend to be efficient, but are highly graph-dependent.

In a more explicit manner, spatial convolutional networks implement convolutions on graphs. The feature representation of each node is iteratively updated by aggregating information from immediate neighbors (Hamilton et al., 2017; Xu et al., 2019) or a receptive field determined by special methods (Niepert et al., 2016). A transformation of the information from neighbors - either linear or non-linear - can be learned through training. The node representation information transferred between vertices are called *messages*. Velicković et al. (2018) used self-attention mechanism to further refine the messages based on local information, namely, features of source and target nodes.

In spatial convolutions, it is essential to have shared filter parameters across different parts of the graph. However, it has been observed that the filters should be adaptive to different local graph structures. In particular, node degrees have been used to reweight the messages or as additional features of node representations (Kipf and Welling, 2017; Monti et al., 2017). This way, messages relevant to hub nodes with high degrees will be different from messages between normal nodes. However, node degree is only the simplest graph structural property. There are much richer and advanced structural information that should be exploited in order to develop structure-adaptive convolutional filters.

Proceedings of the 23rd International Conference on Artificial Intelligence and Statistics (AISTATS) 2020, Palermo, Italy. PMLR: Volume 108. Copyright 2020 by the author(s).

- ▶ Using persistence image to enhance graph neural networks
 - ▶ Beat many GNN models on a broad spectrum of graph benchmarks

► Reweighting via topological information

$$\vec{h}_u^\ell = \sigma \left(\sum_{v \in \overline{\mathcal{N}}(u)} \text{diag}(\tau_{v \rightarrow u}^\ell) W^\ell \vec{h}_v^{\ell-1} \right).$$

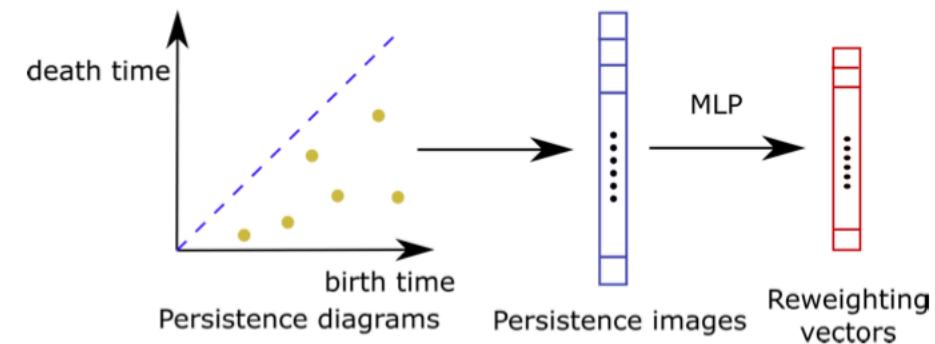
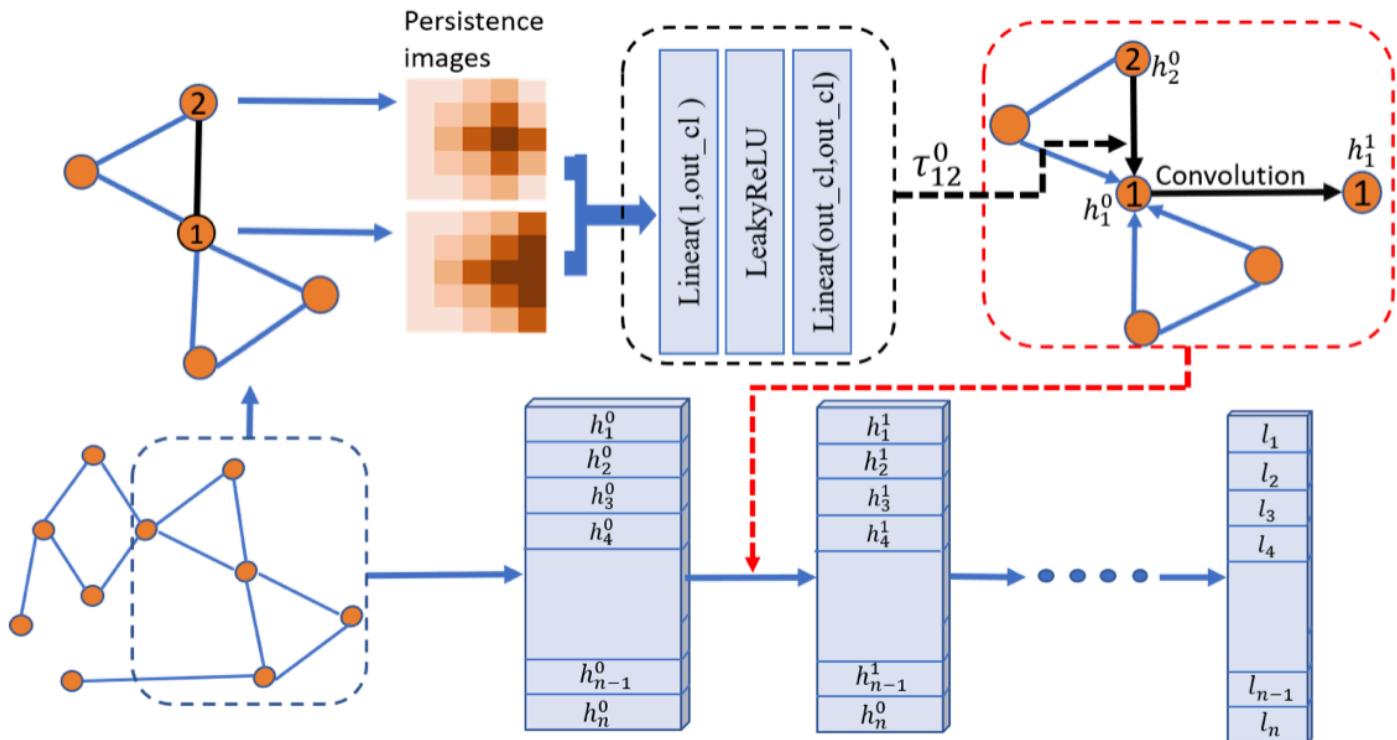


Table 1: Classification Accuracies on Benchmark Datasets

Method	Cora	Citeseer	PubMed	Coauthor CS	Coauthor Physics	Amazon Computer	Amazon Photo
MLP	58.2	59.1	70.0±2.1	88.3±0.7	88.9±1.1	44.9±5.8	69.6±3.8
MoNet	81.7	71.2	78.6±2.3	90.8±0.6	92.5±0.9	83.5±2.2	91.2±1.3
GraphSAGE	79.2	71.2	77.4±2.2	91.3±2.8	93.0±0.8	82.4±1.8	91.4±1.3
U-Net	82.5	72.0	78.9	92.7	94.0	86.0	91.9
WLCN	78.9	67.4	78.1	89.1	90.7	67.6	82.1
GCN	81.5±0.5	70.9±0.5	79.0±0.3	91.1±0.5	92.8±1.0	82.6±2.4	91.2±1.2
GAT	83.0±0.7	72.5±0.7	79.0±0.3	90.5±0.6	92.5±0.9	78.0±19.0	85.1±20.3
PEGN-RC-2	82.7±0.5	71.9±0.6	79.4±0.7	92.9±0.3	94.1±0.3	84.2±1	91.7±0.5
PEGN-RC-1	82.6±0.6	71.7±0.6	78.8±0.5	92.7±0.3	94.2±0.2	86.3±0.6	92.5±0.4

Zhishang Luo¹ Truong Son Hy² Puoya Tabaghi¹ Donghyeon Koh⁴ Michael Defferrard⁴Elahe Rezaei³ Ryan Carey³ Rhett Davis⁵ Rajeev Jain³ Yusu Wang¹¹ University of California San Diego² Indiana State University³ Qualcomm Technologies, Inc.⁴ Qualcomm Wireless GmbH⁵ North Carolina State University

Abstract

The run-time for optimization tools used in chip design has grown with the complexity of designs to the point where it can take several days to go through one design cycle which has become a bottleneck. Designers want fast tools that can quickly give feedback on a design. Using the input and output data of the tools from past designs, one can attempt to build a machine learning model that predicts the outcome of a design in significantly shorter time than running the tool. The accuracy of such models is affected by the representation of the design data, which is usually a netlist that describes the elements of the digital circuit and how they are connected. Graph representations for the netlist together with graph neural networks have been investigated for such models. However, the characteristics of netlists pose several challenges for existing graph learning frameworks, due to the large number of nodes and the importance of long-range interactions between nodes. To address these challenges, we represent the netlist as a directed hypergraph and propose a *Directional Equivariant Hypergraph Neural Network* (DE-HNN) for the effective learning of (directed) hypergraphs. Theoretically, we show that our DE-HNN can universally approximate any node or hyper-edge based function that satisfies certain permutation equivariant and invariant properties natural for directed hypergraphs. We

compare the proposed DE-HNN with several State-of-the-art (SOTA) machine learning models for (hyper)graphs and netlists, and show that the DE-HNN significantly outperforms them in predicting the outcome of optimized place-and-route tools directly from the input netlists. Our source code and the netlists data used are publicly available at <https://github.com/YusuLab/chips.git>

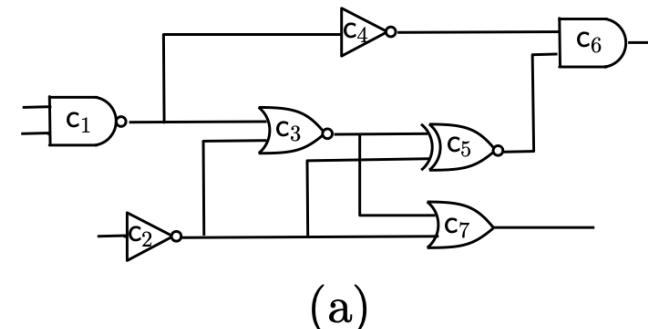
1 Introduction

Chip design is a complicated process involving numerous steps, many of which involve solving hard optimization problems. Just consider the stage of the *place and route* of a synthesized netlist: Here the input is a netlist consisting of *cells* and *nets*, where cells refer to functional units such as logic gates, and nets refer to connections between cells. The goal is to produce a layout of this netlist in a specific 2D region, where gates are placed and connections among them are realized by wires laid out across multiple layers (called “routed”), all while aiming to optimize multiple key properties (e.g., minimizing total wirelength and reducing congested “hotspots”). This place-and-route stage is highly nontrivial to solve for large netlists, and requires a time-consuming process in practice with multiple stages and iterations.

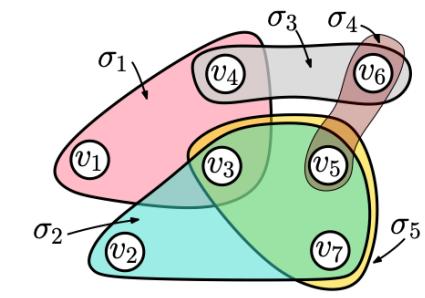
There therefore arises the need for data-driven methods to predict properties of a design directly without the time-consuming place and routing process. To this end, graph neural networks become natural choices, given that the netlists are often represented as a graph or a hypergraph. In this paper, we aim to develop an efficient and effective graph learning architecture to predict post-routing properties (e.g., wirelength or congestion) for a synthesized netlist accurately.

The past decade has witnessed a tremendous growth

► Using persistent homology to create node features of input graphs



(a)



(b)

Table 1: Net-based wirelength regression. Last row “Improvement” refers to the improvement of our full DE-HNN model over the best baseline approach for each metric.

Model	Single-Design			Cross-Design		
	RMSE ↓	MAE ↓	Pearson ↑	RMSE ↓	MAE ↓	Pearson ↑
GCN	1.762	1.276	0.750	1.691	1.276	0.746
GATv2	1.812	1.330	0.687	1.717	1.281	0.737
AllSet	1.718	1.264	0.760	1.837	1.348	0.695
HMPNN	1.841	1.368	0.710	1.785	1.335	0.710
HNHN	1.852	1.368	0.717	1.754	1.333	0.701
NetlistGNN	1.773	1.320	0.740	1.762	1.324	0.718
base DE-HNN	1.751	1.269	0.748	1.731	1.291	0.730
full DE-HNN	1.689	1.245	0.770	1.677	1.242	0.754
Improvement	1.7%	1.6%	1.3%	1.9%	2.6%	1.8%

Table 2: Net-based demand regression for each design.

Model	Single-Design			Cross-Design		
	RMSE ↓	MAE ↓	Pearson ↑	RMSE ↓	MAE ↓	Pearson ↑
GCN	9.321	6.163	0.570	6.571	5.024	0.365
GATv2	9.342	6.118	0.561	6.623	5.137	0.363
AllSet	9.072	5.745	0.632	6.120	4.820	0.345
HMPNN	9.342	6.118	0.561	6.979	5.356	0.306
HNHN	9.342	6.118	0.561	6.390	4.870	0.358
NetlistGNN	9.063	5.839	0.623	8.328	6.839	0.367
base DE-HNN	8.997	5.764	0.630	6.778	5.085	0.337
full DE-HNN	8.381	5.334	0.683	6.037	4.670	0.372
Improvement	7.5%	7.2%	8.1%	1.4%	4.1%	1.4%

Topological loss

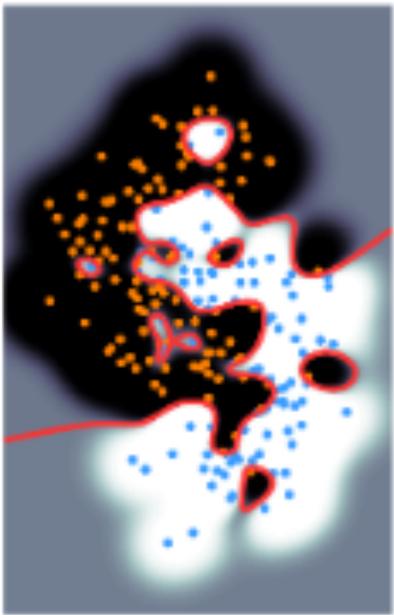
Motivation

- ▶ Many learning / inference problems can be thought of as optimization problems
- ▶ Typical neural network architectures could be “structure-oblivious”
 - ▶ e.g, a CNN is not effective at capturing a global loop feature
- ▶ Hence it makes sense to be able to optimization a function concerning topological properties of domain or data

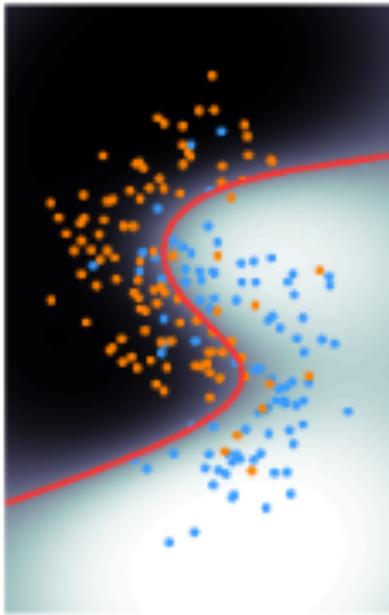
Example 1: Regularizing Classifier Function

- ▶ Classifier function $f: X \rightarrow R$ defined on feature space X
 - ▶ e.g, in binary classification case, $S_f = f^{-1}(0)$ is classification boundary
- ▶ Loss function: $L(f)$
 - ▶ Quality of f in making predictions + **regularization**
- ▶ One issue:
 - ▶ Structure agnostic

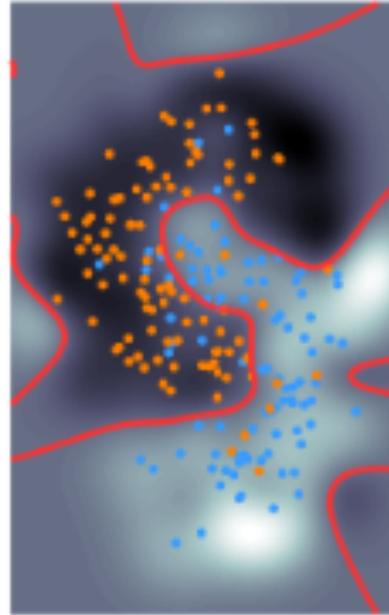
Example 1: Regularizing Classifier Function



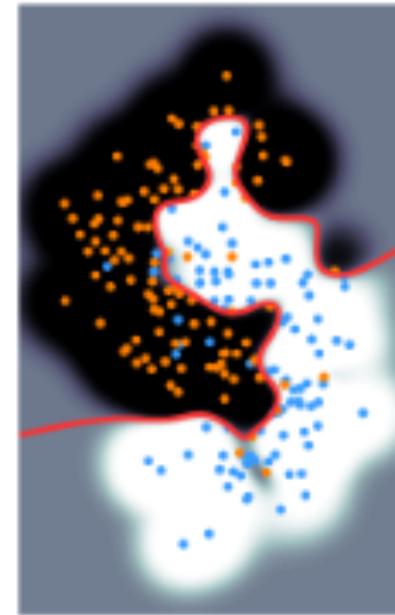
Classical kernel
method; small σ



Classical kernel
method; large σ



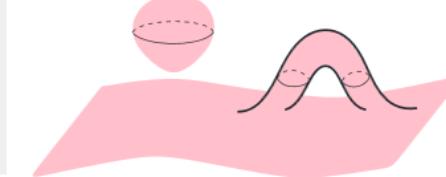
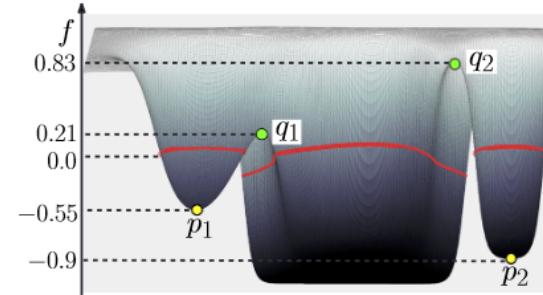
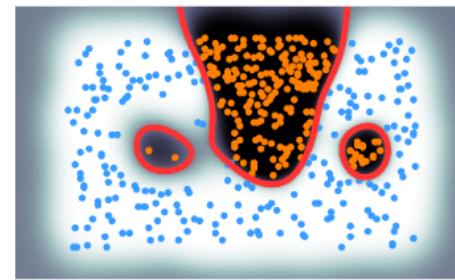
DGR by [Bai et al.,
ICML 2016]



Topological penalty
based loss

Main Idea

- ▶ [Chen, Ni, Bai, Wang, 2019]
- ▶ A topological loss $L_T(f)$

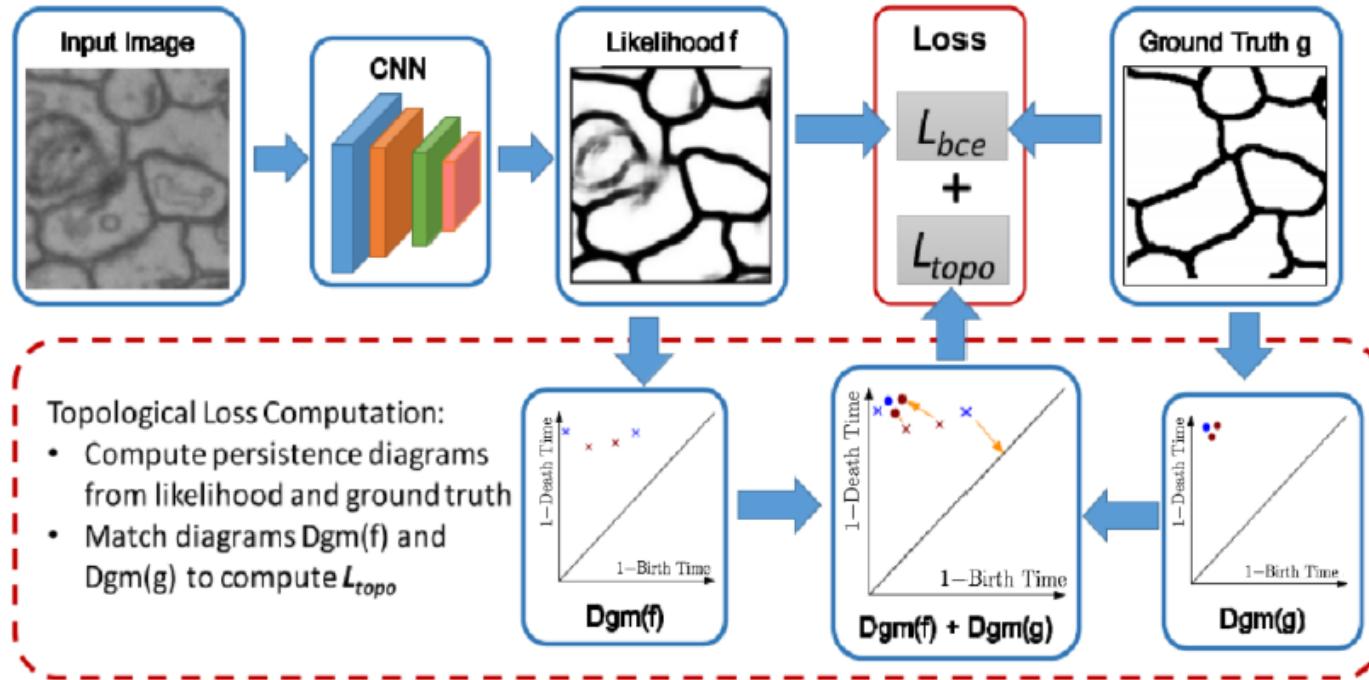


$$L(f, \mathcal{D}) = \sum_{(x,t) \in \mathcal{D}} \ell(f(x, w), t) + \lambda \underbrace{L_T(f(\cdot, w))}_{\text{depends on connected components of the classification boundary}}, \quad (3.1)$$

in which λ is the weight of the topological penalty, L_T . And $\ell(f(x, w), t)$ is the standard per-data loss, e.g., cross-entropy loss, quadratic loss or hinge loss.

depends on connected components of the classification boundary

Example 2: Topological faithfulness of segmentation



Topology-Preserving Deep Image Segmentation

*Xiaoling Hu¹, Li Fuxin², Dimitris Samaras¹ and Chao Chen¹

¹Stony Brook University
²Oregon State University

► [Hu et al., 2019]

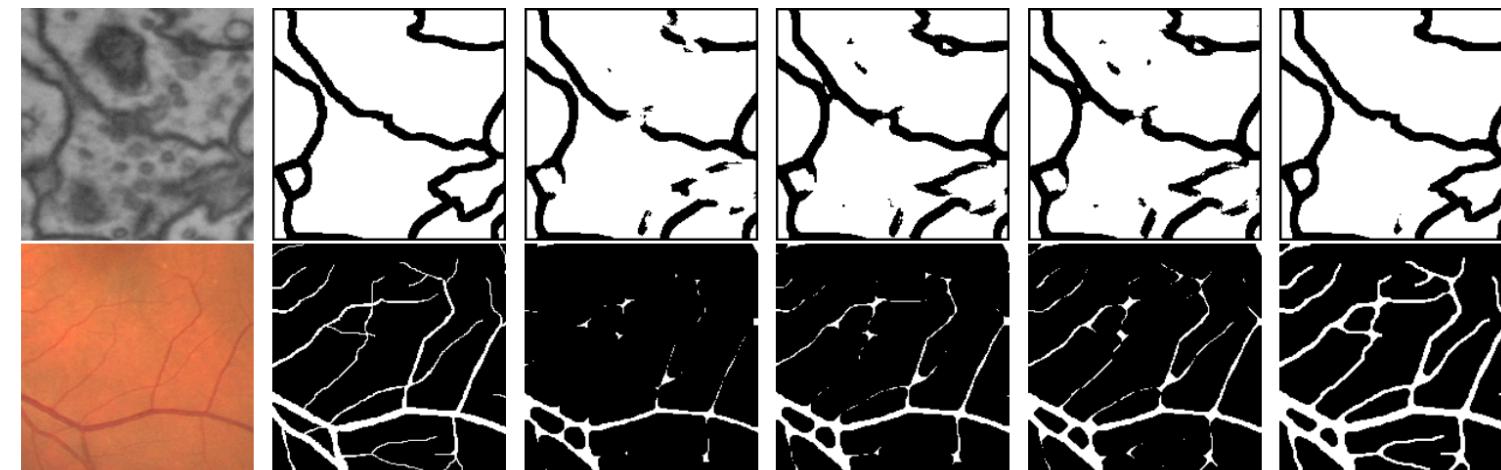
Let Γ be the set of all possible bijections between $Dgm(f)$ and $Dgm(g)$. The loss $L_{topo}(f, g)$ is:

$$\min_{\gamma \in \Gamma} \sum_{p \in Dgm(f)} \|p - \gamma(p)\|^2 = \sum_{p \in Dgm(f)} [\text{birth}(p) - \text{birth}(\gamma^*(p))]^2 + [\text{death}(p) - \text{death}(\gamma^*(p))]^2 \quad (2.2)$$

where γ^* is the optimal matching between two different point sets.

Example 2: Topological faithfulness of segmentation

Dataset	Method	Accuracy	ARI	VOI	Betti Error
DRIVE	DIVE	0.9549 ± 0.0023	0.8407 ± 0.0257	1.936 ± 0.127	3.276 ± 0.642
	U-Net	0.9452 ± 0.0058	0.8343 ± 0.0413	1.975 ± 0.046	3.643 ± 0.536
	Mosin.	0.9543 ± 0.0047	0.8870 ± 0.0386	1.167 ± 0.026	2.784 ± 0.293
	TopoLoss	0.9521 ± 0.0042	0.9024 ± 0.0113	1.083 ± 0.006	1.076 ± 0.265
CrackTree	DIVE	0.9854 ± 0.0052	0.8634 ± 0.0376	1.570 ± 0.078	1.576 ± 0.287
	U-Net	0.9821 ± 0.0097	0.8749 ± 0.0421	1.625 ± 0.104	1.785 ± 0.303
	Mosin.	0.9833 ± 0.0067	0.8897 ± 0.0201	1.113 ± 0.057	1.045 ± 0.214
	TopoLoss	0.9826 ± 0.0084	0.9291 ± 0.0123	0.997 ± 0.011	0.672 ± 0.176
Road	DIVE	0.9734 ± 0.0077	0.8201 ± 0.0128	2.368 ± 0.203	3.598 ± 0.783
	U-Net	0.9786 ± 0.0052	0.8189 ± 0.0097	2.249 ± 0.175	3.439 ± 0.621
	Mosin.	0.9754 ± 0.0043	0.8456 ± 0.0174	1.457 ± 0.096	2.781 ± 0.237
	TopoLoss	0.9728 ± 0.0063	0.8671 ± 0.0068	1.234 ± 0.037	1.275 ± 0.192



sample image, ground truth,

DIVE,

U-Net,

Mosin.,

TopoLoss

Gradient descent of topological loss

- ▶ Challenges:
 - ▶ Can we “differentiate” a function encoding topological information ?
- ▶ Turns out that we can compute gradients for topological loss function based on persistence homology (PH)

Differentiation of a PH-based topological function

- ▶ Suppose we are now given a topological function
 - ▶ $L_T(f) := L_T(D_f)$ that depends on persistent points in a persistence diagram D_f induced by a function f
 - ▶ Now assume the function $f = f_\omega$ is parameterized by parameters ω
 - ▶ Say, f_ω could be modeled by a neural network
- ▶ **Theorem** ([Chen, Ni, Bai, Wang, 2019], [Gameiro, Hiraoka, Obayashi, 2016], [Poulenard, Skraba, Ovsjanikov, 2018])
 - ▶ Using piecewise linear approximation \hat{f}_ω of f_ω , the topological loss $L_T(\hat{f}_\omega)$ is differentiable almost everywhere w.r.t. parameters ω .

How?

- ▶ Suppose we consider $f: X \rightarrow R$,
 - ▶ its persistent diagram $D_f = \left\{ (b_i, d_i), i \in [1, s] \right\}$
- ▶ Recall, each of b_i/ d_i in fact is the function value of some critical point.
 - ▶ Let $\rho : Dgm(f) \rightarrow X \times X$ be such that for any $p_i = (b_i, d_i)$ in $Dgm(f)$, $\rho(p_i)$ is the corresponding pair of critical points.
- ▶ Assume $\omega \in R^m$ is the parameter that parametrize $f_\omega: X \rightarrow R$
 - ▶ Let $\xi: R^m \rightarrow 2^{V \times V}$ be the map $\xi(\omega) = \left\{ (v_{\ell_i}, v_{r_i}) \right\}_{(b_i, d_i) \in D_{f_\omega}} = P_\omega(Dgm(f_\omega))$

How?

Proposition 13.9. Suppose $f_\omega : |K| \rightarrow \mathbb{R}$ is a PL-function with distinct values on all vertices V of K , and K is a finite simplicial complex. Then there exists a neighborhood of ω in the parameter space such that ξ remains constant within this neighborhood; that is, the image set $\xi(\omega) = \rho_\omega(\text{Dgm } f_\omega)$ remains the same for all parameters within this neighborhood.

- ▶ Now given a topological loss function $L_T(f_\omega) := L_T(D_{f_\omega})$
 - ▶ In order to compute $\nabla_w L_T(f_w)$, one needs to compute $\frac{\partial b_i}{\partial w}$ and $\frac{\partial d_i}{\partial w}$
 - ▶ $\frac{\partial b_i}{\partial w} = \frac{\partial f_w(\rho_w(b_i))}{\partial w} = \frac{\partial f_w(v_{\ell_i})}{\partial w}$, which can be computed; and similarly for $\frac{\partial d_i}{\partial w}$

Synthetic							
	KNN	LG	SVM	EE	DGR	KLR	TopoReg
Blob-2 (500,5)	7.61	8.20	7.61	8.41	7.41	7.80	7.20
Moons (500,2)	20.62	20.00	19.80	19.00	19.01	18.83	18.63
Moons (1000,2,Noise 0%)	19.30	19.59	19.89	17.90	19.20	17.80	17.60
Moons (1000,2,Noise 5%)	21.60	19.29	19.59	22.00	22.30	19.00	19.00
Moons (1000,2,Noise 10%)	21.10	19.19	19.89	24.40	26.30	20.00	19.70
Moons (1000,2,Noise 20%)	23.00	19.79	19.40	30.60	30.20	19.50	19.40
AVERAGE	18.87	17.68	17.70	20.39	20.74	21.63	16.92
UCI							
	KNN	LG	SVM	EE	DGR	KLR	TopoReg
SPECT (267,22)	17.57	17.20	18.68	16.38	23.92	18.31	17.54
Congress (435,16)	5.04	4.13	4.59	4.59	4.80	4.12	4.58
Molec. (106,57)	24.54	19.10	19.79	17.25	16.32	19.10	12.62
Cancer (286,9)	29.36	28.65	28.64	28.68	31.42	29.00	28.31
Vertebral (310,6)	15.47	15.46	23.23	17.15	13.56	12.56	12.24
Energy (768,8)	0.78	0.65	0.65	0.91	0.78	0.52	0.52
AVERAGE	15.46	14.20	15.93	14.16	15.13	13.94	11.80
Biomedicine							
	KNN	LG	SVM	EE	DGR	KLR	TopoReg
KIRC (243,166)	30.12	28.87	32.56	31.38	35.50	31.38	26.81
fMRI (1092,19)	46.70	74.91	74.08	82.51	31.32	34.07	33.24

A Topological Regularizer for Classifiers via Persistent Homology

Chao Chen¹

Xiuyan Ni²

Qinxun Bai³

Yusu Wang⁴

¹Stony Brook University, Stony Brook, NY ²City University of New York, New York, NY

³Hikvision Research America, Santa Clara, CA ⁴Ohio State University, Columbus, OH

Other TDA+ML examples

- ▶ Using PH to study activation network
- ▶ Topological complexity of neural network?
- ▶ Using Mapper to understand a trained neural network
- ▶ ...

Journal of Machine Learning Research 21 (2020) 1-40

Submitted 4/20; Revised 7/20; Published 8/20

2021 IEEE International Conference on Big Data (Big Data)

Activation Landscapes as a Topological Summary of Neural Network Performance

Matthew Wheeler
Department of Medicine
University of Florida
Gainesville, USA
mwheeler1@ufl.edu

Jose Bouza
CISE Department
University of Florida
Gainesville, USA
josejbouza@gmail.com

Peter Bubenik
Department of Mathematics
University of Florida
Gainesville, USA
peter.bubenik@ufl.edu

Topology of Deep Neural Networks

Gregory Naitzat
Department of Statistics
University of Chicago
Chicago, IL, 60637, USA

GREGN@UCHICAGO.EDU

Andrey Zhitnikov
Faculty of Electrical Engineering
Technion – Israel Institute of Technology
Haifa 32000, Israel

ANDREYZ@TECHNION.AC.IL

Lek-Heng Lim
Computational and Applied Mathematics Initiative
University of Chicago
Chicago, IL, 60637, USA

LEKHENG@UCHICAGO.EDU